

How I Learned to Stop Worrying and Love the R Console

Irfan Kanat
Department of Information Systems
Arizona State University

November 2, 2015

Outline

- 1 Introduction
- 2 Familiar Examples
- 3 R Console
- 4 Importing Data
- 5 Packages
- 6 Sample Analysis and Visualizations
- 7 Reporting
- 8 Where to Go Next?

Who am I?

Irfan Kanat, PhD Candidate

R user since 2006

Open Source Evangelist

What is this about?

A brief introduction to R.

- R Console
- Importing Data
- Packages
- Sample analyses
- Basic visualization
- Where to get help?
 - Documentation
 - Stack Exchange
 - R community



What is R?

From R project web site:

R is a language *and* an environment for statistical computing and graphics.

- Language
- Environment
- Statistics and Visualization




What is R?

All this means R is very flexible, which played a huge role in its success.

My take: Low cost, high quality, open source solution for your analysis needs.

When to Use R?

R is very strong for your classical machine learning and statistical analysis. Thousands of packages address almost all analysis needs. It is a logical first stop to start analysis.

¹Except when it is not. There are packages to overcome these issues. 

When to Use R?

R is very strong for your classical machine learning and statistical analysis. Thousands of packages address almost all analysis needs. It is a logical first step to start analysis.

Yet it's core design is starting to show its age. There are certain down sides to traditional R:

- Everything is stored in memory¹
- R is single core¹

¹Except when it is not. There are packages to overcome these issues.

Best Part of R

Packages CRAN houses over 7K packages. Providing functionality way beyond what is available in commercial packages.

Community Millions of users mean, all your questions are either already answered or will be in hours.

Performance While memory and core restrictions are real, for the cost of a single user license of a commercial package, you can buy better hardware to run R. Furthermore, with the packages providing multicore and flatfile functionality, R performance is on par or better than commercial packages

Outline

- 1 Introduction
- 2 Familiar Examples**
- 3 R Console
- 4 Importing Data
- 5 Packages
- 6 Sample Analysis and Visualizations
- 7 Reporting
- 8 Where to Go Next?

Logistic Regression

```
# Fit the model
logit_0 <- glm(admit ~ ., admitData, family = "binomial")
# Display fitted model
summary(logit_0)
```

##

Call:

glm(formula = admit ~ ., family = "binomial", data = admitData)

##

Deviance Residuals:

##	Min	1Q	Median	3Q	Max
##	-1.6268	-0.8662	-0.6388	1.1490	2.0790


##

Coefficients:

##		Estimate	Std. Error	z value	Pr(> z)
##	(Intercept)	-3.989979	1.139951	-3.500	0.000465 ***
##	gre	0.002264	0.001094	2.070	0.038465 *
##	gpa	0.804038	0.331819	2.423	0.015388 *
##	rank2	-0.675443	0.316490	-2.134	0.032829 *

Logistic Regression


```
ggplot(admitData, aes(x = gre, y = admit)) + geom_point(aes(colour = rank)) +
  stat_smooth(method = "glm", family = "binomial", , se = T)
```



Support Vector Machine

```
# Fit the model
svm_0 <- svm(admit ~ ., data = admitData, type = "C-classification")
# Plot the results
plot(svm_0, admitData, gre ~ gpa) # Let us plot the results
```

SVM classification plot



Questions



Outline

- 1 Introduction
- 2 Familiar Examples
- 3 R Console**
- 4 Importing Data
- 5 Packages
- 6 Sample Analysis and Visualizations
- 7 Reporting
- 8 Where to Go Next?

Command Driven Interface

Command line may be intimidating

Power or Convenience

Consider the number of

- Functions
- Parameters
- Data sources
- Variables
- Replications


Command Driven Interface

Command line may be intimidating

Power or Convenience

Consider the number of

- Functions
- Parameters
- Data sources
- Variables
- Replications



R Studio

The screenshot displays the R Studio interface with four main panels:

- Script Editor (Top Left):** Contains an R script with comments and code. A red box labeled 'A' highlights the first 44 lines, which include a header, a comment about the workspace, and a series of R console examples.
- Environment (Top Right):** Shows the current environment with variables: `admtData` (400 obs. of 4 variables), `csv` (24 obs. of 3 variables), `d` (10 obs. of 3 variables), `xls` (29 obs. of 3 variables), `values` (a list of 5 elements), `country` (a list of 10 elements), `fac` (a list of 10 elements), `logit_0` (a list of 30 elements), `svm_0` (a list of 30 elements), and `functions` (a list of 30 elements). A blue box labeled 'C' highlights the `values` variable.
- Console (Bottom Left):** Shows the output of the R script. A green box labeled 'B' highlights the output of the `plot(svm_0, admtData, gre-gpa)` command, which displays the SVM classification results.
- Plots (Bottom Right):** Shows an SVM classification plot titled "SVM classification plot". The x-axis is labeled "gpa" and ranges from 2.5 to 4.0. The y-axis is labeled "gre" and ranges from 300 to 800. The plot shows two classes of data points (circles and crosses) separated by a decision boundary. A yellow box labeled 'D' highlights a specific region of the plot.

New Project

File > New Project

Empty Directory > Empty Project > Directory Name: Workshop

R as a Calculator I

```
# Arithmetics
```

```
2 + 2
```

```
## [1] 4
```

```
2 * 3
```

```
## [1] 6
```

```
2^3
```

```
## [1] 8
```

```
log(100, 10)
```

```
## [1] 2
```

R as a Calculator II

```
# Logic
```

```
1 == 2
```

```
## [1] FALSE
```

```
1 != 2
```

```
## [1] TRUE
```

```
2 < 3
```

```
## [1] TRUE
```

Variables I

Variables II

```
A <- 2
```

```
A
```

```
## [1] 2
```

```
a # Case sensitive
```

```
## Error in eval(expr, envir, enclos): object 'a' not found
```

```
"A" != "a" # Explanation
```

```
## [1] TRUE
```

```
B <- 7
```

```
A + B
```

```
## [1] 9
```

Indexes and Data Frames I

```
1:30
```

```
## [1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25 26 27 28 29 30
```

```
C[3]
```

```
## [1] 7
```

```
C[c(2, 3)]
```

```
## [1] 3 7
```

```
C[1:3]
```

```
## [1] 1 3 7
```


Indexes and Data Frames II

```
Countries <- data.frame(names=c("US", "TR", "DE"), supply=c(10, 8, 7),  
                        those=c(TRUE, FALSE, FALSE))
```

Countries

```
##   names supply those  
## 1    US     10  TRUE  
## 2    TR      8 FALSE  
## 3    DE      7 FALSE
```

Indexes and Data Frames III

```
Countries[2, ]
```

```
##      names supply those  
## 2      TR          8 FALSE
```

```
Countries[, 3]
```

```
## [1]  TRUE FALSE FALSE
```

```
Countries[2, 3]
```

```
## [1] FALSE
```

```
Countries[1:2, ]
```

```
##      names supply those  
## 1      US      10  TRUE  
## 2      TR       8 FALSE
```

Loops in R

CAUTION!

R is notoriously inefficient with your classic loops

- Structure of the Data Frame
- Memory Management

Try to use an apply function instead.

Vectorize your operations.

For Loop in R

```
for (i in 1:3) print(i)
```

```
## [1] 1
```

```
## [1] 2
```

```
## [1] 3
```

```
# Iterating through a data frame
```

```
for (i in 1:nrow(Countries)) {  
  print(Countries[i, ])  
}
```

```
## names supply those
```

```
## 1 US 10 TRUE
```

```
## names supply those
```

```
## 2 TR 8 FALSE
```

```
## names supply those
```

```
## 3 DE 7 FALSE
```

Functions I

```
ls() # List the contents of the environment
```

```
## [1] "A"          "admitData" "B"          "c"          "C"  
## [6] "Countries" "HelloWorld" "i"          "logit_0"    "logit_1"  
## [11] "svm_0"
```

```
mean(C) # Takes parameters
```

```
## [1] 5
```

```
mean(C, trim = 0.1, na.rm = T) # Takes multiple parameters
```

```
## [1] 5
```

```
log(sum(C)/length(C)) # Can be combined
```

```
## [1] 1.609438
```

Functions II

```
HelloWorld <- function(x, y = 1) {  
  for (i in 1:y) {  
    print(paste("Hello", x))  
  }  
}
```

```
HelloWorld("MSBA")
```

```
## [1] "Hello MSBA"
```

```
HelloWorld("MSBA", 2)
```

```
## [1] "Hello MSBA"
```

```
## [1] "Hello MSBA"
```

Functions III

```
HelloWorld  # Review the source code
```

```
## function(x, y = 1) {
##   for (i in 1:y) {
##     print(paste("Hello", x))
##   }
## }
```

```
ls
```

```
## function (name, pos = -1L, envir = as.environment(pos), all.names = FALSE,
##   pattern, sorted = TRUE)
## {
##   if (!missing(name)) {
##     pos <- tryCatch(name, error = function(e) e)
##     if (inherits(pos, "error")) {
##       name <- substitute(name)
##       if (!is.character(name))
##         name <- deparse(name)
##       warning(gettextf("%s converted to character string",
##
```

Commonly Used Functions I

```
ls() # Get a list of objects in the workspace
```

```
## [1] "A"          "admitData"  "B"          "c"          "C"
## [6] "Countries"  "HelloWorld" "i"          "logit_0"    "logit_1"
## [11] "svm_0"
```

```
ls(pattern = "*_0") # partial match on object search
```

```
## [1] "logit_0" "svm_0"
```

```
rm("svm_0") # Remove an object from the workspace
```

```
# rm(list=ls(pattern=ls())) # This would remove everything if ran
```


Commonly Used Functions II

```
mean(A)  # Mean
```

```
## [1] 2
```

```
sd(admitData[, "gre"])  # Standard Deviation
```

```
## [1] 115.5165
```

```
AIC(logit_0)
```

```
## [1] 470.5175
```

Commonly Used Functions III

```
str(Countries) # Look at the structure of objects
```

```
## 'data.frame': 3 obs. of 3 variables:
## $ names : Factor w/ 3 levels "DE","TR","US": 3 2 1
## $ supply: num 10 8 7
## $ those : logi TRUE FALSE FALSE
```

```
summary(Countries) # Get summary of data
```

```
## names      supply      those
## DE:1  Min.   : 7.000  Mode :logical
## TR:1  1st Qu.: 7.500  FALSE:2
## US:1  Median : 8.000  TRUE :1
##      Mean   : 8.333  NA's :0
##      3rd Qu.: 9.000
##      Max.   :10.000
```

Commonly Used Functions IV

```
psych::describe(admitData)  # Better version of summary stats in psych pack
```

##	vars	n	mean	sd	median	trimmed	mad	min	max	range	sk
## admit	1	400	0.32	0.47	0.0	0.27	0.00	0.00	1	1.00	0.
## gre	2	400	587.70	115.52	580.0	589.06	118.61	220.00	800	580.00	-0.
## gpa	3	400	3.39	0.38	3.4	3.40	0.40	2.26	4	1.74	-0.
## rank*	4	400	2.48	0.94	2.0	2.48	1.48	1.00	4	3.00	0.

##	kurtosis	se
## admit	-1.39	0.02
## gre	-0.36	5.78
## gpa	-0.60	0.02
## rank*	-0.91	0.05

Commonly Used Functions V

```
summary(logit_1) # Get summary of model

##
## Call:
## glm(formula = admit ~ gre, family = "binomial", data = admitData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.1623  -0.9052  -0.7547   1.3486   1.9879
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.901344   0.606038  -4.787 1.69e-06 ***
## gre          0.003582   0.000986   3.633 0.00028 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 499.98  on 399  degrees of freedom
## Residual deviance: 486.06  on 398  degrees of freedom
## AIC: 490.06
##
## Number of Fisher Scoring iterations: 4
```


Commonly Used Functions VI

```
cor(admitData[, 1:3]) # Get correlation matrix
```

```
##           admit           gre           gpa
## admit 1.0000000 0.1844343 0.1782123
## gre    0.1844343 1.0000000 0.3842659
## gpa    0.1782123 0.3842659 1.0000000
```

```
pairs(admitData) # Visualize correlation table
```

Commonly Used Functions VII



Questions



Outline

- 1 Introduction
- 2 Familiar Examples
- 3 R Console
- 4 Importing Data**
- 5 Packages
- 6 Sample Analysis and Visualizations
- 7 Reporting
- 8 Where to Go Next?

Importing Data

R allows importing data from a wide variety of sources.

- Comma Separated Values (CSV)
- Databases
- Flat files
- Lesser statistical packages
- and more

Importing CSV Files

CSV has certain advantages that make it popular.

- Compatibility
- Flexibility
- Simplicity

Sample

```
"iso2", "Supply", "Those"
"AU", 20, 0
"TR", 80, 1
"US", 100, 0
"GB", 50, 0
"DE", 70, 0
```

We use `read.csv()` or `read.csv2()` commands to import the csv files.

```
saveData <- read.csv("PathToCSV", header = TRUE, sep = ",", quote = "\"", )
```

Working with Excel Files

Much like CSV, except it lacks the simplicity, flexibility, and compatibility of CSV.

```
# Load the necessary library  
library(xlsx)  
# Read in the data from excel file  
xlsx <- read.xlsx("country.xlsx", sheetIndex = 1)
```

Working with Databases

No speed advantage.

Data larger than memory.

Working with databases:

- Work in the database.
- Import data from database.



Working with Databases

```
# Load the necessary library
library(RMySQL)
# Establish connection to the database.
channel <- dbConnect(MySQL(), user = "uname", password = "pwd", host = "127.0.0.1",
  dbname = "exampledata")
# Send query and save results in R workspace
sql <- dbGetQuery(channel, "SELECT * FROM table;")
```

Lesser Statistical Packages :P

Foreign Package

Newer file formats

- sas7bdat
- readstata13



Questions



Outline

- 1 Introduction
- 2 Familiar Examples
- 3 R Console
- 4 Importing Data
- 5 Packages**
- 6 Sample Analysis and Visualizations
- 7 Reporting
- 8 Where to Go Next?

Packages: Source of R's Power

Encountered already

Make R extendible

Like libraries

Collection of:

- functions
- documentation
- data files



Gifts from the Community

Currently over 7000 packages

for

- Statistical Modeling
- Machine Learning
- Data Manipulation
- Visualization
- ...

from

- Economics
- Computer Science
- Statistics
- Medicine
- ...



Great but Where are My Gifts?

Comprehensive R Archive Network (CRAN)

A Group of FTP and HTML servers hosting R packages.

R has built in package management facilities.

Package Management

```
# Installing a package
install.packages("e1071") # Notice the quotes around package name
# Loading package into memory
library(e1071) # Notice the lack of quotes
# Unload package
detach("package:e1071", unload = TRUE) # Notice the package: prepended
```

```
# Get the list of packages loaded
(.packages())
```

```
## [1] "e1071"      "ggplot2"    "knitr"      "stats"      "graphics"
## [6] "grDevices" "utils"      "datasets"   "methods"    "base"
```

```
# Get list of all installed packages (part of output omitted)
.packages(all.available = T)
```

```
## [1] "acepack"      "assertthat"  "bdsmatrix"   "betareg"
## [5] "BH"           "biglm"       "bit"         "bitops"
## [9] "brew"        "car"         "caTools"     "chron"
```

How to Find Packages

If you want to search a certain word in installed packages' documentation, you can always use `??` or `help.search()`

```
??mixed  
help.search("mixed model")
```

Internet searches are a bit problematic as R can be a bit ambiguous until Google learns you are interested in the statistical computing environment.

Comprehensive R Archive Network (CRAN)

R Forge

R site search also available with command `RSiteSearch()`

R seek

Commonly Used Packages: Data Manipulation

data.tables Replaces traditional `data.frame`.

- Faster access/write
- Improved selection
- Improved subsetting
- Improved aggregation

Not a drop-in replacement as it breaks compatibility in some cases.

ddplyr

Additional functionality for:

- selection
- filtering
- aggregation

Provides efficient back-end data structures to speed things up.

Works with databases as well.

Commonly Used Packages: Statistics

Multivariate Regression: Base package

Generalized Linear Models: glm package

Traditional Econometric Models: plm package

Mixed Modeling: nlme and lme4 packages

Commonly Used Packages: Machine Learning

Classifiers (KNN, LVQ): class package

Support Vector Machines: kernlab, e1071 packages

Clustering: Base package (kmeans(), hclust()), mclust package

Neural Networks: neuralnet package.

Questions



Outline

- 1 Introduction
- 2 Familiar Examples
- 3 R Console
- 4 Importing Data
- 5 Packages
- 6 Sample Analysis and Visualizations**
- 7 Reporting
- 8 Where to Go Next?

Questions



Outline

- 1 Introduction
- 2 Familiar Examples
- 3 R Console
- 4 Importing Data
- 5 Packages
- 6 Sample Analysis and Visualizations
- 7 Reporting**
- 8 Where to Go Next?

Questions



Outline

- 1 Introduction
- 2 Familiar Examples
- 3 R Console
- 4 Importing Data
- 5 Packages
- 6 Sample Analysis and Visualizations
- 7 Reporting
- 8 Where to Go Next?**